

PRODUCT REQUIREMENTS DOCUMENT

AI-assisted portfolio building and hosting, designed for work that deserves context.

A robust, privacy-conscious platform that turns resumes, work artifacts, links, and media into polished, editable portfolios and hosted professional presence.

DOCUMENT TYPE

Detailed PRD

RELEASE TARGET

MVP through V1

PREPARED

June 29, 2026

Prepared for product, design, engineering, QA, security, and operations alignment.

Document purpose and reading guide

This document defines the product, technical, design, quality, privacy, and operating requirements for YourWork. It is written to align a cross-functional team before implementation and to provide concrete criteria for shipping a secure MVP.

Product thesis

Professional portfolios are difficult to assemble because a person's best evidence is spread across resumes, project links, presentations, screenshots, videos, credentials, and informal narratives. YourWork helps people organize that evidence, generate a compelling first draft, and retain meaningful editorial control.

Operating principle

AI accelerates authorship; it does not obscure it. Every generated section should be traceable to user-supplied source material, editable before publication, and clearly governed by user privacy and publication settings.

How to interpret this PRD

LABEL	MEANING IN THIS DOCUMENT
MUST	Required to ship the phase safely or meet the intended product promise.
SHOULD	Important for quality, but may be deferred with an explicit product decision.
REVIEW GATE	Requires deliberate engineering, security, legal, design, or operations review before broad release.
OUT OF SCOPE	Not planned for MVP; included to prevent accidental expansion.

Table of contents

01. Executive summary	3	07. Backend, infrastructure, and data model	11
02. Users and jobs to be done	4	08. GitHub repos and delivery workflow	14
03. Goals, success metrics, scope	5	09. Security, privacy, and trust	15
04. MVP experience and functional requirements	6	10. Quality, QA, and release gates	17
05. AI assistance system	8	11. Phased delivery roadmap	19
06. Design system and frontend	10	12. Risks, decisions, and acceptance criteria	22

YourWork PRD - Internal product artifact. Product choices in this document are proposals until validated through discovery, technical spikes, legal review, and user research.

YourWork gives people a credible digital home for their work.

YourWork is an AI-assisted portfolio-building and hosting platform. Users upload or connect their professional materials, choose a template or permit AI-directed layout, review a draft portfolio, refine it in an editor, and publish it to a hosted URL or connected custom domain.

Input

Resume, work document, project links, media, bio, profile image, credentials, and goals

Assist

Structured extraction, content synthesis, information architecture, draft copy, page composition, and alt text suggestions

Publish

Responsive hosted page, privacy controls, SEO controls, shareable URL, and later custom domain support

Problem statement

Most portfolio tools ask users to start with a blank canvas or a generic template. That approach creates high cognitive load, produces incomplete portfolios, and makes it hard for people whose work is not naturally represented by a single image gallery. Conversely, fully automated generation risks inaccurate claims, generic language, and loss of author agency.

YourWork's differentiation: it treats the portfolio as an evidence-backed narrative. Users supply source material; AI turns material into structured, editable candidate content; templates translate that content into a coherent public experience; the user controls what is private, public, and deleted.

Primary product promise

AUDIENCE PROMISE	PLATFORM RESPONSE	PROOF OF QUALITY
I have work scattered across many places.	Upload, link, and organize materials in one workspace.	Import completes, assets are categorized, and source provenance is visible.
I do not know how to present myself.	AI proposes narrative, sections, layout, and concise case-study drafts.	User can accept, revise, reject, or regenerate at section level.
I need a portfolio I can share confidently.	Responsive hosted site with privacy and publication controls.	Published page meets accessibility, performance, and link-validation targets.

Non-goals for MVP

OUT OF SCOPE

Full social network

OUT OF SCOPE

Native mobile apps

OUT OF SCOPE

Public marketplace

OUT OF SCOPE

Team collaboration

OUT OF SCOPE

Paid subscriptions

Target users, needs, and experience principles

1. Emerging professional

Student, career changer, junior creative, early engineer, or self-taught specialist.

Need: Turn partial experience into a credible presentation without overclaiming.

Success: Publishes a portfolio with at least three evidence-backed sections and can share it in an application.

2. Experienced operator

Designer, engineer, marketer, manager, consultant, researcher, or founder with extensive artifacts.

Need: Distill a large body of work into focused case studies while controlling confidential details.

Success: Produces a tailored site, controls case-study visibility, and can update it without a developer.

3. Freelancer or creator

Independent professional who needs an elegant web presence and inbound lead pathway.

Need: Show services, projects, testimonials, and contact information in a brand-aligned place.

Success: Uses a public URL as the primary destination for prospective clients.

4. Career support professional

Mentor, educator, workforce program staff, or coach assisting an individual.

Need: Help someone produce a portfolio while respecting ownership and consent.

Success: Can guide a user through source selection and review without taking control of the user's account.

Core jobs to be done

SITUATION	JOB	PRODUCT IMPLICATION
When I am preparing for an opportunity...	Help me turn my real experience into a concise story that shows evidence and outcomes.	Guided import, content inventory, AI drafting, case study sections, and user approval.
When I have confidential or sensitive work...	Help me show my capabilities without exposing information I cannot publish.	Private drafts, per-project visibility, redaction warnings, no automatic public publishing.
When I want a site but do not want to design one...	Help me choose a layout that feels professional and still reflects me.	Template gallery, AI template recommendation, brand controls, and reversible choices.
When my work changes...	Help me update the relevant part of my portfolio without rebuilding everything.	Structured content model, modular blocks, revision history, and publish diff.

Design principle: reduce blank-page anxiety without replacing judgment. The system should ask for just enough context to create a useful draft, then make editorial decisions legible and reversible.

What MVP must prove

The MVP should prove that users can move from scattered materials to a trustworthy, visually polished, published portfolio without requiring professional web-design or engineering skills.

Goals

GOAL	MVP MEASURE	TARGET DIRECTION
Fast time to first useful draft	Median time from first asset upload to generated preview	Under 10 minutes for a typical starter portfolio
Meaningful completion	Users who publish after generating a draft	Measure baseline in beta; optimize with research before setting hard target
User control and trust	Published changes that were reviewed in editor vs. published directly after generation	High review participation; no default automatic publication
Portfolio quality	Accessibility, valid links, responsive rendering, and completion score	All launch-critical checks pass before publish
Operational reliability	Generation jobs, publish jobs, and page delivery success	Set SLOs before beta; instrument from day one

Scope summary

MVP included

- Email/password and OAuth authentication.
- Workspace, profile, resume/work-document upload, project links, media uploads, and profile image.
- AI content extraction, structured profile/project drafts, template recommendation, and portfolio generation.
- Template-driven visual editor with manual edits and AI-assisted edits.
- Public hosting at a platform subdomain; publish, unpublish, preview, and basic SEO controls.
- Asset storage, content moderation boundaries, privacy settings, deletion, audit trail, and analytics events.


V1 or later

- Custom domains, billing and subscriptions, advanced analytics, collaboration, multilingual generation, and marketplace templates.
- Multi-page portfolio sites, lead CRM, e-commerce, public commenting, and native mobile apps.
- Deep integrations beyond import links such as cloud-drive sync, GitHub profile sync, or LinkedIn import.

Constraints and assumptions

CONSTRAINT	DECISION
Users may upload personal and professional material.	Use private-by-default storage and never make imported material public without explicit publish action.
LLM output can be inaccurate or stylistically generic.	Require source grounding, content provenance, user review, and per-section regeneration.
Some uploads may be large, unsupported, or unsafe.	Validate type and size, scan uploads, create async processing, and show recoverable failure states.
Launch resources are limited.	Use managed services for auth, storage, databases, observability, and hosting rather than bespoke infrastructure.

The end-to-end user journey



Functional requirement matrix

AREA	REQUIREMENT	PRIORITY	ACCEPTANCE SIGNAL
Account creation	Support sign-up with email/password and OAuth. Verify email for publish-capable accounts.	MUST	User can create, sign in, reset password, and terminate sessions.
Onboarding	Capture role, portfolio objective, primary audience, experience level, chosen mode: template-first or AI-directed.	MUST	Inputs populate portfolio brief and can be edited later.
Source intake	Accept PDF, DOCX, TXT, common image formats, video links, URLs, and manual entry. Validate before processing.	MUST	Supported files upload securely and receive meaningful processing status.
Project records	Create editable project entities with title, role, timeframe, context, contribution, outcomes, media, links, and privacy level.	MUST	User can create, edit, reorder, hide, and delete projects.
Templates	Offer a curated gallery of at least 4 accessible responsive templates with preview and style settings.	MUST	Template choice preserves content and can be switched without loss.
AI generation	Generate structured content and page composition from selected sources, with citations to source fragments internally.	MUST	Draft includes a reviewable set of sections, never auto-publishes, and retains source links.
Editing	Provide block-level editing for copy, images, sections, order, visibility, and design tokens. Provide AI rewrite with instruction field.	MUST	Every change appears in preview and is persisted safely.
Publishing	Support draft, preview, publish, unpublish, and republish. Create canonical platform URL and robots/SEO controls.	MUST	Public site reflects the selected revision and updates with a clear publish state.
Privacy controls	Set whole-site and per-project visibility: draft-only, unlisted, public. Publish consent is explicit.	MUST	Private or unlisted content is not indexable or publicly discoverable.
Contact	Allow optional contact button using email reveal or protected form in later phase.	SHOULD	User can enable/disable contact details without exposing unintended data.

Deliberate-review features: authentication, uploads, public publishing, custom domains, payments, analytics, and contact forms must not be treated as routine UI tasks. Each requires security, privacy, and abuse review before release.

AI should create a useful draft, not an unaccountable persona.

YourWork's AI agent operates as a constrained content and layout assistant. It receives only user-approved source material, produces structured candidate fields and page blocks, shows what it used, and keeps the user in control of publication.

AI capabilities for MVP


Content intelligence

- Extract basic entities from uploaded resume or work document: employers, titles, dates, skills, projects, education, and credentials.
- Classify assets and links into potential project evidence.
- Suggest project titles, summaries, role statements, impact statements, and section hierarchy.
- Identify missing information and ask focused follow-up questions.
- Generate accessible image alt-text drafts and concise page metadata.

Design intelligence

- Recommend a template based on content density, role, and audience.
- Map structured content to supported template blocks.
- Recommend ordering, media placement, and emphasis.
- Respect template constraints rather than producing unmaintainable one-off HTML.
- When "AI-directed" is chosen, select among curated design decisions, not unrestricted visual generation.

AI generation pipeline



Grounding and quality rules

RULE	IMPLEMENTATION REQUIREMENT
Evidence sensitivity	Maintain source references per generated claim when technically feasible; surface a "review sources" affordance in the editor.
No invented outcomes	Prompts must instruct the model not to fabricate metrics, employers, credentials, dates, client names, or project outcomes. Ambiguity becomes a follow-up question or neutral wording.
User editorial control	Allow section-level accept, edit, regenerate, dismiss, and restore. Store user edits separately from generated draft state.
Content boundaries	Do not use sensitive user material to train product models unless a separate, explicit opt-in is designed and approved.
Observability without unnecessary retention	Log job metadata, versions, failure categories, and redacted diagnostics. Minimize raw prompt/content retention.

AI review gate: Before beta, test for unsupported claims, inappropriate tone, privacy leakage from mixed-visibility sources, prompt injection embedded in documents, and unstable rendering from malformed generated content.

Frontend architecture: create without trapping users in a black box.

Recommended frontend stack

LAYER	RECOMMENDATION	RATIONALE
Framework	Next.js with TypeScript	Supports fast product UI, server rendering for public sites, route-level security patterns, and static/edge delivery choices.
Component system	Accessible internal design system on top of vetted primitive components	Ensures consistent inputs, dialogs, menus, editor controls, and accessible behavior.
Styling	Token-based CSS variables with utility or component-scoped styling	Enables template themes, user brand settings, dark/light considerations, and maintainable visual consistency.
State and data	Server-first data fetching; query cache only for editor interactions	Reduces stale state while preserving responsive editing and preview behavior.
Editor model	Schema-driven block editor, not freeform HTML	Preserves security, template portability, accessibility, and predictable rendering.
Validation	Shared schema library for API payloads and editor content	Prevents divergent assumptions between frontend and backend.

Key user surfaces

Workspace

Portfolio list, progress checklist, generation status, recent changes, publish state, and recovery actions.

Content studio

Source inventory, profile fields, project records, asset organizer, privacy labels, and missing-information prompts.

Portfolio editor

Desktop/tablet/mobile preview, template switcher, section tree, block controls, AI assistance, history, and publish review.

Template system requirements

- Templates must be code-owned, versioned, responsive, keyboard navigable, and compatible with the same content block schema.
- MVP templates: Minimal Professional, Editorial Case Study, Visual Showcase, and Builder/Technical Portfolio.
- Every template must support a coherent default hierarchy: hero, about, selected work, project detail, experience/skills, and contact/links.
- Template customization must use bounded controls: accent color, type scale, density, visual emphasis, navigation style, and section visibility.
- Generated content should never bypass the block schema or inject arbitrary styles/scripts.


Accessibility and performance requirements

AREA	REQUIREMENT
Accessibility	Keyboard-operable editor, visible focus states, semantic structure, accessible labels, valid heading order, contrast checks, alt-text pathway, and reduced-motion respect.
Responsive behavior	Preview must represent desktop, tablet, and mobile; public pages must not rely on the editor's desktop layout.
Performance	Responsive images, CDN delivery, lazy loading below the fold, video embeds deferred, and page budgets defined before launch.
SEO	User-editable title and description, canonical URLs, structured metadata where appropriate, robots controls, sitemap eligibility for public pages, and no indexing for unlisted pages.

Backend architecture: managed foundations, explicit boundaries

The platform should separate interactive user-facing APIs from long-running ingestion and generation work. It should favor managed services for operationally sensitive capabilities while retaining a clear domain model in the application backend.

Recommended service topology



Backend stack proposal

CAPABILITY	RECOMMENDATION	IMPLEMENTATION NOTES
Application APIs	Python FastAPI or TypeScript NestJS service, with one primary backend language selected before Phase 1	Both fit well. Favor the team's stronger operational expertise; share API contracts and schemas with frontend.
Database	PostgreSQL	Use relational integrity for users, organizations later, portfolio content, revisions, permissions, jobs, and audit events.
Queue and workers	Managed queue plus idempotent workers	Separate file processing, link metadata extraction, AI generation, image transforms, and publish compilation.
Object storage	Private bucket with signed upload/download URLs and CDN for public optimized assets only	Never expose original private uploads by predictable URL.
Search / retrieval	Postgres full-text plus vector extension only if source grounding needs semantic retrieval at scale	Start simple; justify vector retrieval with observed relevance needs.
Auth	Managed identity provider or vetted auth service	Use secure password storage, OAuth, email verification, sessions, MFA readiness, and rate protection.
Observability	Centralized logs, metrics, traces, error tracking, and product analytics	Redact sensitive content by default; establish alert thresholds before beta.

Core API domains

Interactive APIs

Identity, account settings, portfolios, profile, projects, assets, templates, editor blocks, revisions, publish states, privacy settings, signed uploads, and user notifications.

Async job APIs

Document ingestion, virus scan status, OCR/text extraction where permitted, media optimization, external-link metadata fetch, AI generation, quality checks, static publish compilation, and cleanup/deletion jobs.

A schema that preserves ownership, revisions, and traceability

Use a normalized core model with JSONB only where flexible template-specific properties are genuinely necessary. Every publishable page should be reproducible from a recorded revision.

TABLE	KEY FIELDS	PURPOSE AND CONSTRAINTS
users	id, email, display_name, avatar_asset_id, status, created_at	Primary account. Email is unique and normalized. Status supports active, suspended, and deleted lifecycle.
auth_identities	id, user_id, provider, provider_subject, verified_at	Maps local account to password/OAuth provider identities. Provider subject unique per provider.
workspaces	id, owner_user_id, name, plan, created_at	Future-proofing for multiple portfolios or collaboration without exposing multi-user features in MVP.
workspace_members	workspace_id, user_id, role, invited_at	Keep schema ready for future collaboration; MVP role is owner only.
portfolios	id, workspace_id, slug, title, template_id, status, visibility, active_revision_id	Logical portfolio. Slug unique at platform scope. Status includes draft, published, archived.
portfolio_revisions	id, portfolio_id, revision_number, created_by, source, content_snapshot, template_version, created_at	Immutable publishable snapshot. Source records user/manual, AI, system, or migration.
portfolio_blocks	id, revision_id, block_type, sort_order, content_json, visibility	Structured blocks: hero, text, project_collection, project_detail, gallery, experience, links, contact.
profiles	id, workspace_id, full_name, headline, bio, location_text, primary_role, public_fields	Structured identity fields. Store only what user provides; separate private metadata from public fields.
projects	id, workspace_id, title, role, start_date, end_date, summary, confidentiality_level, status	Canonical project content referenced by blocks. Confidentiality supports private, redact, public review labels.
project_contributions	id, project_id, contribution_type, description, outcome_text, sort_order	Models contribution and result claims independently for editing and provenance.
assets	id, workspace_id, storage_key, media_type, mime_type, byte_size, checksum, scan_status, created_at	Object metadata only. Storage key is non-guessable. Scan status gates use in public pages.
asset_variants	id, asset_id, variant_kind, width, height, storage_key	Optimized derivatives, thumbnails, transcodes, and responsive images.
source_documents	id, workspace_id, asset_id, doc_type, extraction_status, text_encrypted_ref	Represents resume/work-document inputs. Extracted text handling must follow retention policy.
external_links	id, workspace_id, url, canonical_url, title, metadata_json, fetch_status	Links to projects or profiles; SSRF-safe fetching and controlled metadata extraction.
source_fragments	id, source_document_id, locator, text, confidence, sensitivity_label	Optional ground-truth fragments for traceability. Encrypt or minimize when containing sensitive content.
generation_jobs	id, workspace_id, portfolio_id, job_type, status, input_manifest, model_version, prompt_version, output_ref	Asynchronous AI/document jobs. Idempotency key prevents duplication; do not store more raw content than necessary.
generated_claims	id, generation_job_id, entity_type, entity_id, claim_text, confidence, review_status	Enables source-backed review of generated statements and user acceptance/rejection metrics.
content_provenance	id, target_type, target_id, source_fragment_id, relation_type	Links generated or edited content to source evidence where available.
publication_events	id, portfolio_id, revision_id, action, actor_user_id, occurred_at, metadata	Publish/unpublish/revert audit trail and operational debugging.
audit_events	id, user_id, workspace_id, action, target_type, target_id, ip_hash, created_at	Security-relevant and privacy-relevant event history with retention controls.
consents		

	id, user_id, consent_type, policy_version, granted_at, revoked_at	Records explicit consent for terms, analytics preference, and any future training opt-in.
custom_domains	id, portfolio_id, hostname, status, verification_token, certificate_status	V1 feature: exact ownership verification and lifecycle tracking.

Data integrity and lifecycle rules

- Use soft-delete and retention timers for user-deletable content; purge underlying object storage according to deletion policy.
- Ensure asset access is checked against workspace membership and publication visibility on every delivery path.
- Use immutable revision snapshots for published content; never mutate a historical published revision in place.
- Version templates and prompt policies so a published revision remains reproducible enough to debug or roll back.
- Use foreign keys, uniqueness constraints, check constraints for enum-like values, and idempotency keys for async workflows.

Repository strategy and engineering operating model

Start with a modular monorepo only if the team is comfortable enforcing boundaries; otherwise use a small number of focused repositories with shared contracts. The objective is fast collaboration without coupling deploys or leaking secrets.

Recommended repository topology

yourwork-web
Next.js application: authenticated workspace, content studio, editor, public rendering, template packages, visual regression tests, web analytics events.

yourwork-api
Primary API service: domain models, authorization, portfolio CRUD, asset signed URLs, publishing API, integrations, OpenAPI contract, unit/integration tests.

yourwork-workers
Async workers: document ingestion, safe metadata fetch, asset processing, AI orchestration, publish compilation, deletion/retention jobs.

yourwork-infra
Infrastructure as code, environment modules, secrets references, networking, storage, database, queues, monitors, dashboards, and runbooks.

yourwork-contracts
Shared schemas/types: API request and response schemas, portfolio block specification, events, error taxonomy, fixture content, and generated client code.

Branching and CI/CD standards

PRACTICE	REQUIREMENT
Protected main branch	Require pull request review, passing unit/type/lint tests, secret scanning, dependency scanning, and successful build before merge.
Environment flow	Local -> preview/ephemeral environment -> staging -> production. Production promotion requires tagged release, migration review, and rollback plan.
Database migrations	Forward-only migrations with dry-run in staging, backup/rollback plan, and deploy ordering documented for each migration.
Infrastructure changes	Plan output reviewed in pull request; no manual console drift for production resources except documented incident response.
Secrets	Never in source control, logs, test snapshots, or front-end bundles. Use managed secret storage and environment-scoped access.
Release notes	Maintain user-facing and operator-facing notes. Flag migrations, privacy changes, model/prompt changes, and data-retention changes.

Definition of done for an engineering story

- Acceptance criteria implemented and demonstrable.
- Unit tests and relevant integration tests pass.
- Accessibility and responsive review complete for UI changes.
- Logging, error handling, and analytics event behavior defined.

- Security/privacy impact assessed.
- API schemas updated and versioned where applicable.
- Feature flag or rollback plan exists for risky changes.
- QA test cases updated; documentation/runbook revised where needed.

Trust is a product feature, not a policy page.

YourWork handles professional identity, work samples, media, and potentially confidential documents. Privacy-by-default, secure processing, clear consent, and abuse resilience are launch requirements.

Security baseline

CONTROL AREA	REQUIREMENTS	REVIEW GATE
Authentication	Password hashing through vetted auth provider, email verification, secure session management, OAuth state/PKCE, login rate limiting, account recovery controls, and MFA-ready architecture.	REQUIRED BEFORE BETA
Authorization	Server-side authorization on every resource. Workspace and portfolio ownership checks. Never rely on client-side visibility state alone.	REQUIRED BEFORE BETA
File uploads	Allowlist MIME types and size limits, signed upload URLs, malware scanning, content-type verification, isolated processing, image/video transform safety, and no direct execution.	REQUIRED BEFORE PUBLIC UPLOAD
External URLs	Defend against SSRF and unsafe redirects: DNS/IP filtering, fetch allow/deny policy, rate limits, timeout, content-size limit, and isolated metadata-fetch worker.	REQUIRED BEFORE LINK IMPORT
Public content	Output encoding, CSP, sanitization, no arbitrary script embeds in MVP, safe iframe policies for supported media, anti-enumeration for unlisted content.	REQUIRED BEFORE PUBLISH
Secrets and keys	Managed vault, rotation procedure, least privilege, separate keys per environment, and provider access audit.	REQUIRED BEFORE PRODUCTION
Security monitoring	Audit events, suspicious sign-in monitoring, upload scan alerts, rate-limit telemetry, error alerting, vulnerability and dependency management.	SHOULD BEFORE BETA

Privacy and data governance

User-facing commitments

- All portfolios begin as private drafts.
- Users choose public, unlisted, or draft-only publication state at portfolio and project level.
- Users can download their content and request deletion.
- AI generation is opt-in per action, with a clear description of source material included.
- Publication requires explicit confirmation after a preview review.

Internal governance controls

- Data inventory and retention schedule for source documents, extracted text, assets, logs, and model interactions.
- Data processing agreements and subprocessors reviewed before production.
- Access controls for support and operations; support access only when necessary and auditable.
- Separate product analytics from raw content; use consent-aware event collection.
- Incident response and breach notification procedures documented.

Abuse and content safety

Provide clear policies for impersonation, copyright infringement, harassment, malware distribution, and deceptive portfolio claims. Implement reporting mechanisms for public pages and a moderation escalation path. The MVP may use manual review for edge cases, but its data model and audit tooling should support a future moderation workflow.

Legal and compliance review: Terms of Service, Privacy Notice, cookie/analytics disclosures, user content license terms, deletion mechanics, DMCA/copyright pathway where applicable, and domain policy must be reviewed by qualified counsel for the intended jurisdictions before launch.

Quality strategy: test the system, the generated content, and the public experience.

Test pyramid

LAYER	COVERAGE EXAMPLES	OWNERS
Unit tests	Schema validation, authorization helpers, content transformation, prompt input builders, slug generation, visibility rules, template block validation, utility functions.	Engineers
Component tests	Editor blocks, upload UI, accessibility states, validation messages, template switch behavior, preview controls.	Frontend engineers
API integration tests	Auth flows, signed asset URL issuance, portfolio CRUD, revision creation, publish/unpublish, deletion workflow, rate limits, permission boundaries.	Backend engineers
Worker integration tests	Queue idempotency, malware scan result states, file-type validation, job retries, AI response schema validation, failure recovery, cleanup jobs.	Platform/backend engineers
End-to-end tests	Signup -> upload -> generate -> edit -> preview -> publish -> unpublish -> delete using seeded fixtures and test accounts.	QA and engineering
Visual regression tests	Public template pages across breakpoints, critical editor views, image galleries, long content, missing media, dark/high-contrast scenarios if supported.	Design engineering / frontend
Security tests	Authorization bypass attempts, upload bypass, unsafe URL fetch, XSS, CSRF/session behavior, rate-limit behavior, dependency scans, secrets scans.	Security owner + engineering
AI evaluation tests	Grounding, unsupported-claim rate, tone, prompt-injection resistance, privacy separation, schema validity, deterministic render checks from fixed fixtures.	AI/product/QA

QA test scenarios for MVP

Happy paths

- New user creates an account, imports a resume and two project links, generates a draft, edits a section, and publishes.
- User switches templates after editing and loses no content.
- User can use a platform URL on mobile, desktop, and tablet with functional links and appropriate image loading.
- User unpublishes and verifies public page is no longer accessible/indexable as designed.

Failure and edge paths

- Unsupported/oversized/unsafe files produce clear errors without partial exposure.
- AI job times out or returns malformed output; user can retry or continue manually.
- Mixed-visibility project sources do not leak into a public generated section.
- Concurrent edits, retrying a publish job, and browser refresh during upload do not corrupt revision state.

Release quality gates

GATE	REQUIRED EVIDENCE
Feature complete	Acceptance tests pass; documented known limitations; analytics events and support messaging included.
Beta ready	Authentication, authorization, uploads, private/public visibility, and critical user journey tested in staging; accessibility and mobile review complete.
Public launch ready	Threat-model review, incident response plan, legal/privacy review, performance baseline, monitoring/alerts, backup restore test, support escalation and rollback plan complete.

Build the foundations first, then prove the core outcome.

Time estimates are intentionally omitted because delivery capacity, team composition, vendor selection, and research findings determine sequencing. Each phase exits only when its specified release gate is met.

Phase 0 - Product and technical foundation

Purpose: Align the product contract before implementation.

Deliverables: user research plan, design principles, information architecture, content schema, block schema, template concepts, architecture decision records, threat model, data inventory, analytics plan, API contract outline, and test strategy.

Exit gate: approved MVP scope; validated prototype; documented decisions on backend language, auth vendor, storage, AI provider gateway, and hosting model.

Phase 1 - Backend/frontend infrastructure and developer platform

Purpose: Establish the secure, testable foundation required by all user-facing capabilities.

Deliverables: repositories and CI/CD, local developer environment, staging/production infrastructure, database migrations, auth integration, session management, authorization middleware, API scaffolding, object storage integration, signed uploads, queue/worker skeleton, observability, secret management, audit logging, feature flags, and baseline test harnesses.

Exit gate: authenticated user can create a private workspace and portfolio; all data paths are protected; deploy and rollback are demonstrable; core tests run in CI.

Phase 2 - Content studio and portfolio data model

Purpose: Let users reliably enter, organize, and control their source information before AI generation.

Deliverables: profile editor, project CRUD, asset manager, source-document records, external links, upload validation/scanning states, privacy levels, source inventory, revision model, and manual portfolio section composition.

Exit gate: user can create a complete manual portfolio draft with at least profile, projects, media, and visibility controls; no AI dependency is required to progress.

Phase 3 - Template system, editor, preview, and publishing

Purpose: Deliver a polished, reliable public portfolio experience.

Deliverables: curated templates, design tokens, block editor, live preview, responsive preview modes, accessible public renderer, platform subdomain URL, draft/preview/publish/unpublish workflow, SEO controls, static or dynamic public site delivery, and publish audit trail.

Exit gate: user can manually produce and publish a responsive, accessible portfolio from approved content.

Phase 4 - AI-assisted content and composition

Purpose: Reduce creation time while retaining accuracy, ownership, and transparency.

Deliverables: document extraction, AI generation jobs, source manifests, structured profile/project suggestions, template recommendation, section drafting, AI-directed layout selection from templates, section-level rewrite, generation status UX, provenance indicators, and evaluation fixtures.

Exit gate: selected beta users generate usable drafts from representative inputs with reviewable output, controlled failures, no automatic public publishing, and monitored quality.

Phase 5 - Beta hardening and controlled launch

Purpose: Validate user value and operational readiness with a manageable audience.

Deliverables: beta onboarding, support tooling, content reporting workflow, performance optimization, cross-browser QA, accessibility audit, security validation, backup/restore drill, monitoring dashboards, product instrumentation, feedback loops, and launch documentation.

Exit gate: critical journey and security gates pass; beta learnings are triaged; support and incident processes are staffed.

● Phase 6 - V1 expansion

Purpose: Extend from strong MVP to a scalable hosted portfolio business.

Candidate deliverables: custom domains, billing, subscriptions, advanced templates, more granular analytics, lead/contact workflows, multilingual support, integrations, collaboration, and template marketplace foundations.

Exit gate: each expansion feature passes its own security, privacy, operational, and business-model review.

Open decisions and launch-critical risks

Key product and technical decisions

DECISION	RECOMMENDED DIRECTION	WHY IT MATTERS	OWNER
Single vs. dual backend language	Choose one primary service language in Phase 0; use TypeScript end-to-end if speed of shared types dominates, Python if document/AI pipeline expertise is strongest.	Reduces operational complexity and duplicated tooling.	Engineering lead
Public site rendering	Use server-rendered or statically compiled portfolio pages from immutable published revisions.	Improves performance, SEO, reliability, and rollback behavior.	Frontend/platform
AI provider abstraction	Implement a provider gateway with versioned prompt templates, request policies, and evaluation fixtures.	Supports safety controls, cost management, provider flexibility, and debuggability.	AI/platform
Generated content provenance	Start with source-manifest and section-level provenance; evolve to claim-level support where valuable.	Builds trust without overcomplicating MVP.	Product/AI
Custom domains	Defer until the platform URL, publish lifecycle, and security baseline are proven.	Introduces DNS verification, certificate lifecycle, phishing/abuse, support, and billing complexity.	Product/platform

Risk register

RISK	IMPACT	MITIGATION	TRIGGER FOR ESCALATION
AI drafts contain false or inflated claims	Trust harm, reputational harm	Grounded prompts, source-aware review, user approval, neutral wording when uncertain, evaluation set	Repeated qualitative beta feedback or measurable unsupported-claim incidents
Private material appears in public output	High privacy and legal risk	Default private, per-asset/per-project visibility, source manifest filters, publish preview, regression tests	Any confirmed exposure; stop affected publishing path and incident-review
Unbounded media costs	Business risk and degraded performance	File caps, image transforms, quotas, lifecycle rules, CDN caching, roadmap for paid storage tiers	Storage/egress exceeds budget threshold or performance budget
Template system becomes too flexible	Slow development, fragile public pages	Schema-based blocks and bounded design tokens; no arbitrary user code in MVP	Frequent one-off template exceptions or excessive rendering bugs
Auth or upload vulnerability	Account/data compromise	Managed auth, secure storage, scan pipeline, threat modeling, pentest/specialist review before launch	Critical vulnerability or failed security gate

MVP acceptance criteria

Product acceptance

- A new user can create an account, choose a portfolio goal, upload supported sources, and receive clear processing feedback.
- A user can create a portfolio using either a chosen template or an AI-recommended template path.
- AI output is editable and is never published without an explicit user action.
- A user can publish, unpublish, and restore an earlier revision without losing draft content.
- Published portfolios work across current desktop and mobile browsers supported by the QA plan.

Platform acceptance

- Authorization tests demonstrate that users cannot access another user's private assets, portfolios, revisions, or source documents.
- Upload and URL import security controls have been tested against defined threat cases.
- Critical path monitoring, error tracking, backups, and recovery procedures are operational.
- Accessibility, privacy, and legal review gates are documented with explicit outcomes.
- AI evaluation fixtures cover representative resumes, project types, media mixes, missing-data states, and sensitive-content boundaries.

Final product framing: YourWork should feel like a capable design and writing partner, but its durable value comes from a dependable system of evidence, user control, privacy, publishing quality, and maintainable infrastructure.

End of PRD. Recommended next artifact: a Phase 0 product discovery brief with prototype flows, content-schema workshop, threat model, and a prioritized engineering backlog.